



Autumn 2023

Week 3

go.osu.edu/aiclubsignup



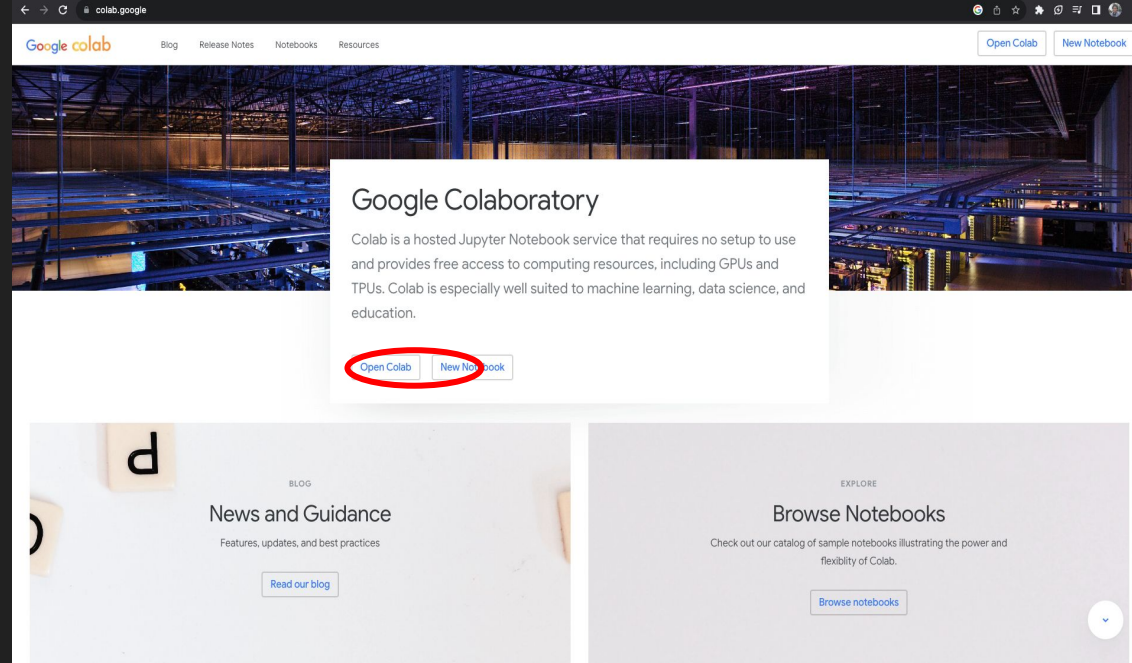
Plan for today

- We're going to explain how to find a dataset
- How to import that dataset into your Google Colab environment
- How to use the built-in datasets functionality of PyTorch
- How to transform your data so that you can feed it into a model
- How to feed a deep learning model the data



Before we start...

- Open up Google Colab
- Select “New Notebook”
- Name the notebook whatever you want
- <https://colab.research.google.com/>



Enter this into a code block

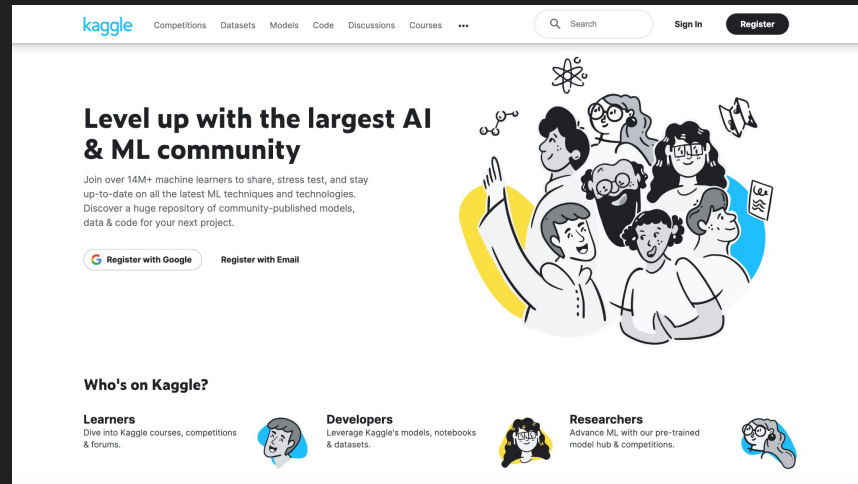
```
[ ] import torch  
    import torchvision  
    from torchvision import transforms
```

- This is installing libraries that we will need for today
 - PyTorch and torchvision (an accompanying library of PyTorch)



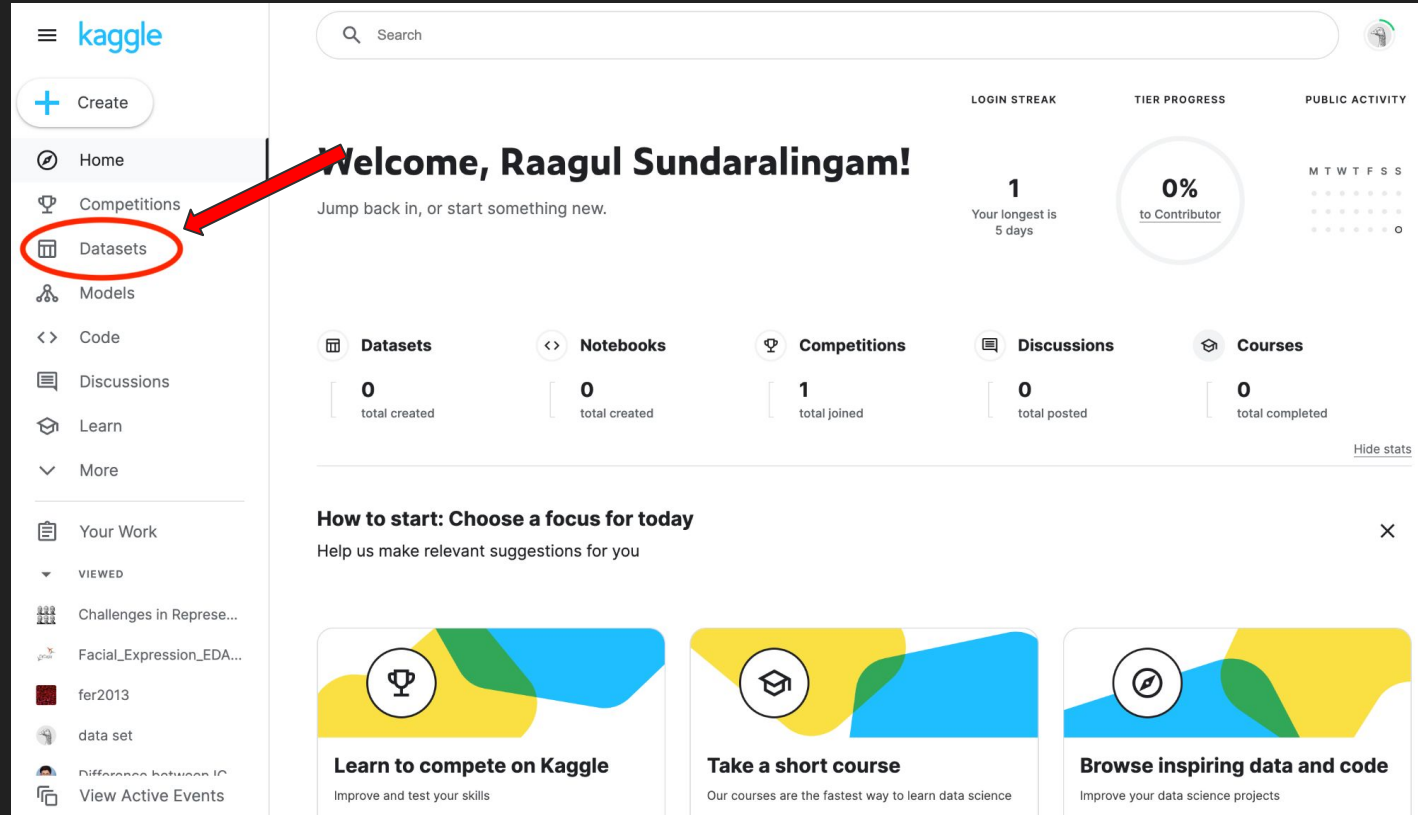
How to find a dataset for your project? (pt. 1)

- Kaggle is a popular machine learning community website that offers tons of datasets, pretrained models, tutorials, and etc.
- Here is the link: <https://www.kaggle.com/>
- If you don't already have an account, you can make one for free!



How to find a dataset for your project? (pt. 2)

Click on the
“Datasets”
option at the left
hand side



The screenshot shows the Kaggle homepage for user Raagul Sundaralingam. The left sidebar contains a navigation menu with the following items: Home, Competitions, Datasets (highlighted with a red circle and a red arrow), Models, Code, Discussions, Learn, More, Your Work, and a list of viewed items including Challenges in Representation, Facial Expression EDA, fer2013, data set, Difference between IC, and View Active Events. The main content area features a welcome message, a search bar, and statistics for Datasets (0 total created), Notebooks (0 total created), Competitions (1 total joined), Discussions (0 total posted), and Courses (0 total completed). Below this, there is a section titled 'How to start: Choose a focus for today' with three cards: 'Learn to compete on Kaggle', 'Take a short course', and 'Browse inspiring data and code'.

Welcome, Raagul Sundaralingam!
Jump back in, or start something new.

Statistics:


Datasets	Notebooks	Competitions	Discussions	Courses
0 total created	0 total created	1 total joined	0 total posted	0 total completed

How to start: Choose a focus for today
Help us make relevant suggestions for you

- Learn to compete on Kaggle**
Improve and test your skills
- Take a short course**
Our courses are the fastest way to learn data science
- Browse inspiring data and code**
Improve your data science projects

How to find a dataset for your project? (pt. 3)

Once you reach this screen, simply click on the search bar on the top and type in whatever topic you want data on (ex: housing prices)



The screenshot shows the Kaggle Datasets page. On the left is a sidebar with navigation links: Create, Home, Competitions, Datasets (highlighted), Models, Code, Discussions, Learn, More, Your Work, and a 'VIEWED' section with recent datasets. The main content area has a top search bar, a 'Datasets' header, a description, and buttons for 'New Dataset' and 'Your Work'. A red arrow points to the search bar. Below is another search bar labeled 'Search datasets' with a 'Filters' button. A row of category tags includes 'All datasets', 'Computer Science', 'Education', 'Classification', 'Computer Vision', 'NLP', 'Data Visualization', and 'Pre-Trained Model'. The 'Trending Datasets' section displays four dataset cards: 'Happiness Report 2022', 'Age, Weight, Height, BMI Analysis', 'Hourly Earnings of Female and Male...', and 'Restaurant Reviews'. Each card shows a thumbnail, title, author, update time, usability score, and file details.

Kaggle

+ Create

Home

Competitions

Datasets

Models

<> Code

Discussions

Learn

More

Your Work

VIEWED

Challenges in Represe...

Facial_Expression_EDA...

fer2013

data set

Difference between IC

View Active Events

Search

Datasets

Explore, analyze, and share quality data. [Learn more](#) about data types, creating, and collaborating.

+ New Dataset Your Work


Search datasets

Filters

All datasets Computer Science Education Classification Computer Vision NLP Data Visualization Pre-Trained Model

Trending Datasets

See All




Happiness Report 2022

Sujay Kapadnis · Updated 12 hours ...

Usability **10.0** · 337 kB

2 Files (other, CSV)




Age, Weight, Height, BMI Analysis

Ruken Missonnier · Updated 9 days ...

Usability **8.8** · 11 kB

1 File (CSV)




Hourly Earnings of Female and Male...

efimpolianskii · Updated 23 days ago

Usability **10.0** · 13 kB

2 Files (CSV)



Restaurant Reviews

Joakim Arvidsson · Updated 7 days ...

Usability **10.0** · 1 MB

1 File (CSV)

Dataset for emotion recognition

- For the emotion recognition project, we will be using the FER2013 dataset.
- There's plenty of datasets out there for this project (a lot of them are better too), but we chose FER2013 because the PyTorch Datasets class already has support for FER2013. **(More on this later...)**
- Here is the link to the official FER2013 dataset on Kaggle: <https://www.kaggle.com/c/challenges-in-representation-learning-facial-expression-recognition-challenge>

The screenshot shows the Kaggle challenge page for the 'Challenges in Representation Learning: Facial Expression Recognition Challenge'. The page has a dark blue header with the challenge title and a '\$500 Prize Money' badge. Below the header, there are tabs for 'Overview', 'Data', 'Code', 'Discussion', 'Leaderboard', 'Rules', and 'Team'. The 'Overview' tab is selected. The main content area is divided into three sections: 'Overview', 'Prizes & Awards', and 'Participation'. The 'Overview' section includes a timeline showing the challenge started on April 12, 2013, and will close on May 24, 2013. The 'Prizes & Awards' section lists '\$500 - Medals and Points'. The 'Participation' section shows '63 Competitors', '56 Teams', and '190 Entries'. A 'Table of Contents' sidebar on the right lists 'Description', 'Evaluation', 'Timeline', 'Prizes', and 'Citation'. The 'Description' section is expanded, showing the challenge's purpose: to design features better and faster than humans can, using a facial expression classification dataset.

Search

Research Prediction Competition

Challenges in Representation Learning: Facial Expression Recognition Challenge

Learn facial expressions from an image

\$500 Prize Money

56 teams · 10 years ago

Overview Data Code Discussion Leaderboard Rules Team

Overview

Start
Apr 12, 2013

Close
May 24, 2013

Prizes & Awards

\$500 · Medals and Points

Participation

63 Competitors
56 Teams
190 Entries

Table of Contents

- Description
- Evaluation
- Timeline
- Prizes
- Citation

Description

One motivation for representation learning is that learning algorithms can design features better and faster than humans can. To this end, we hold this challenge that does not explicitly require that entries use representation learning. Rather, we introduce an entirely new dataset and invite competitors from all related communities to solve it. The dataset for this challenge is a facial expression classification dataset that we have assembled from the internet. Because this is a newly introduced dataset, this contest will see which methods are the easiest to get quickly working on new data.


Example baseline submissions are available as part of the `pyearn2` python package available


Downloading the FER2013 Dataset into Google Colab

- Click on the “Data” option under the banner of the FER2013 link
- Scroll down and click on the “Download All” option (should be in a black button)
- It should’ve downloaded as a zip file called “challenges-in-representation-learning-facial-expression-recognition-challenge.zip”



Downloading the FER2013 Dataset into Google Colab

- Click on the  icon on the left side bar.
- Then right click in the file area and click “Upload” (as shown in the picture to the right)
- Then upload the zip file that we downloaded in the previous slide



Downloading the FER2013 Dataset into Google Colab

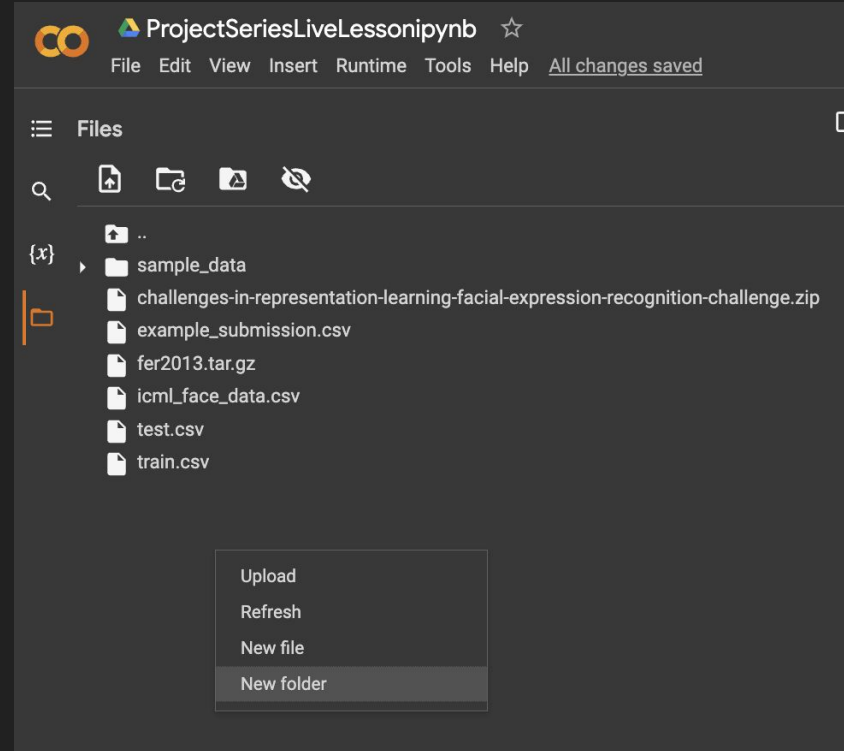
- Create a new code block with the button that says “ + Code ” at the top.
- Type in this line into that code block and click run (if your zip file is called something else, replace the “challenges-in-representation-learning-facial-expression-recognition-challenge” part with the name of your zip file)
- Then wait a minute or two for the files to show up on your colab files.



```
!unzip /content/challenges-in-representation-learning-facial-expression-recognition-challenge.zip
```

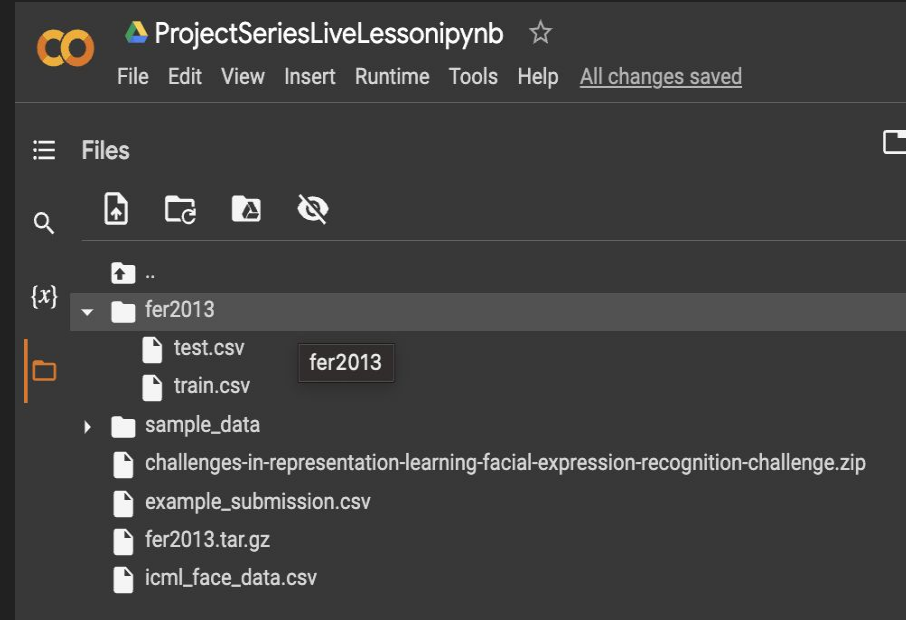
Downloading the FER2013 Dataset into Google Colab

- Once all the files unzipped, right click in the file area and click “New folder” (as shown in the picture to the right)
- Title this folder “fer2013” (MAKE SURE TO COPY PASTE THIS EXACT)



Downloading the FER2013 Dataset into Google Colab

- Move the “train.csv” and “test.csv” files into the “fer2013” folder you just made.
- You should have something that looks like the picture to the right.



We Have the Dataset in Your Google Colab

You did it! You now have the dataset set up in your Google Colab environment.

- Unfortunately you will have to do this setup every time your notebook loses connection (which it automatically does if you are inactive).
- SO DON'T DELETE THE ZIP FILE ON YOUR COMPUTER
- It will come in handy when you have to reupload the dataset back into your notebook.



Now to the Good Stuff

```
def train_pl():  
    #the transformation we will apply to the images from the FER2013 dataset  
    transform = transforms.Compose([  
        transforms.Grayscale(),  
        transforms.ToTensor(), # Convert image to tensor  
        transforms.Normalize(0.485, 0.229) # Normalize image  
    ])  
  
    # loading the data from the directory I have stored the downloaded FER2013 dataset  
    train_data = torchvision.datasets.FER2013(root='/content', split = 'train', transform=transform)  
  
    # create dataloaders so that the FER2013 data can be loaded into the model we will implement  
    train_loader = torch.utils.data.DataLoader(train_data, batch_size=19, shuffle=True, num_workers=2)  
  
    return train_loader
```

PyTorch Datasets

```
# loading the data from the directory I have stored the downloaded FER2013 dataset  
train_data = torchvision.datasets.FER2013(root='/content', split = 'train', transform=transform)
```

There's multiple parts to this statement:


- torchvision.datasets.FER2013
- 1st parameter: root = '/content'
- 2nd parameter: split = 'train'
- 3rd parameter: transform = transform

Let's look at each part



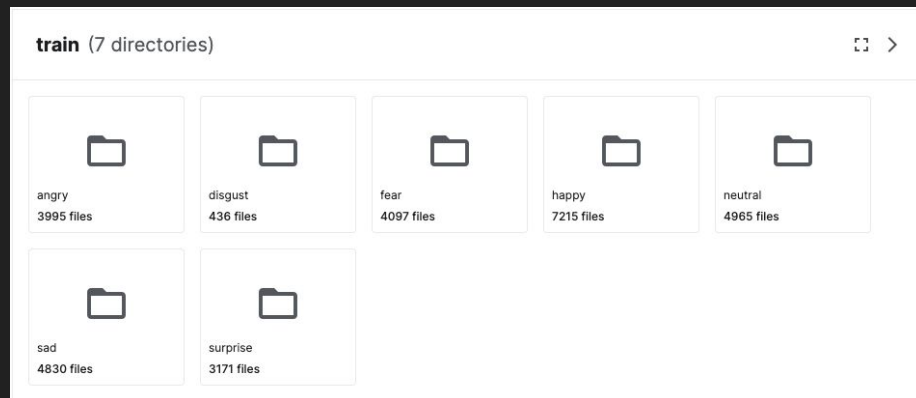
FER2013

- FER2013 is a highly popular dataset for facial emotion recognition
- It was released back in 2013 on Kaggle as a research attempt to open up the improvement of facial detection and recognition algorithms to the public.
- Still widely used today to train facial emotion recognition systems.



FER2013 (cont.)

- Datasets have labels and the corresponding data
- 7 classifications:
 - Happy, neutral, sad, fear, angry, surprise, and disgust
- 48 x 48 pixel images
- Grayscale images
 - A pixel can have a value between 0 (black) and 255 (white)
- Unequal distribution (look to the image on the right)



PyTorch Datasets

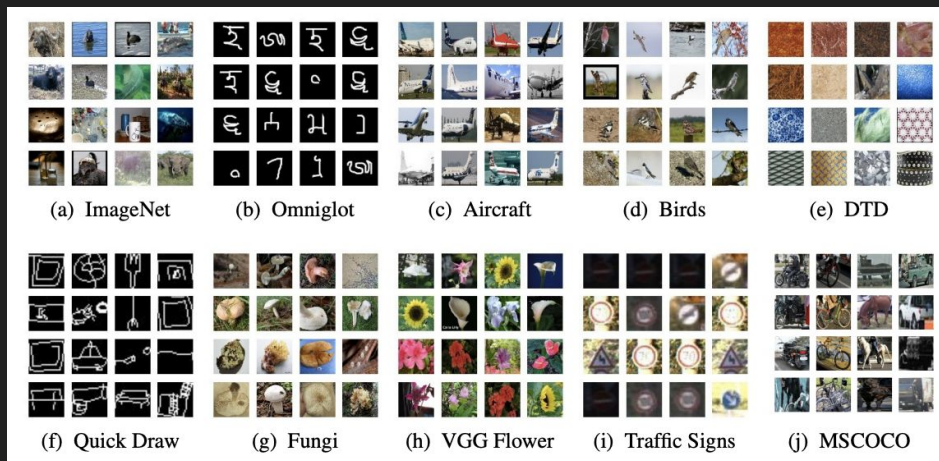
Dataset: a class in PyTorch that defines a way to represent your dataset so that it is organized, structured, and easy to use with PyTorch.

It allows you to:

- Get the length of your dataset
- Get individual data sample in your dataset

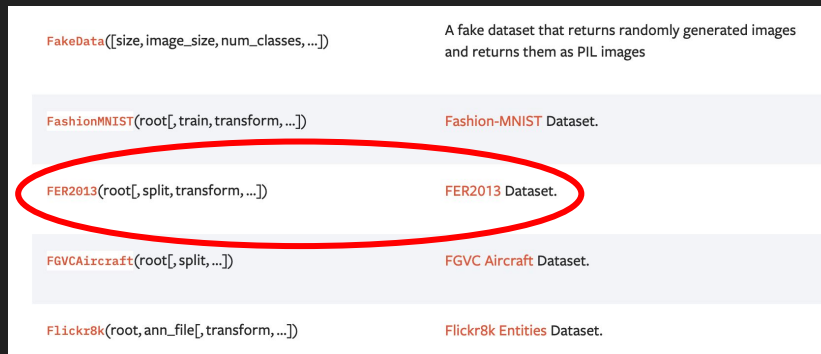
You could either:

- Use the built-in datasets provided by PyTorch (what we will be doing)
- Create your own custom PyTorch dataset by extending the Dataset class (tutorial can be found [here](#) at the part that says “Creating a Custom Dataset for your files”)




PyTorch Datasets

- Remember when we said that the PyTorch Datasets class already has built-in support for FER2013?
- If you visit the documentation for the [Pytorch Datasets class](#), you'll see a list of very popular datasets for different deep learning use cases that PyTorch has built in support. One of those datasets is FER2013 ...
- You should find it under the “Image Classification” section




torchvision.datasets.FER2013

- “torchvision.datasets.FER2013” is how you access the built-in PyTorch dataset definition for FER2013 (for beginners, think of how you use System.out.print() in Java).
- “torchvision” is PyTorch’s accompanying library.
- “datasets” is a specific module in the library
- “FER2013” is the code for the class within “datasets” that defines how to use FER2013 with PyTorch
- Info on Python Modules can be found [here](#)




1st parameter: root = '/content'

- The root parameter should be set to a file path
- A file path is a set of directions the computer uses to navigate your file system and find what it needs
- In this case, the fer2013 folder we set up can be found in content directory of your notebook.
 - Kind of like Pomerene 160 can be found in Pomerene Hall



2nd parameter: split = 'train'

- We're telling torchvision that we want the training data from the dataset you can find at the file path we gave it.
- Datasets are usually split into 3 sections: training, validation, and testing.
- Best practice:
 - Training: 70%
 - Testing: 20%
 - Validation: 10%
- We won't be going over validation in this course, we will only be going over training and testing.



3rd parameter: transform = transform

- Before we start feeding the model our data, we need to adjust it with the specifications of our dataset. Make it more digestible and useful for our model.
- “transforms.Compose” is a function in torchvision that allows us to make a series of transformations.
- We want to convert our data samples to grayscale.
- We want to convert each sample to a tensor (required for feeding data into a model, more on this next week)
- We want to normalize our data so that we can standardize each sample
 - This involves statistics we won't be going over this course, but the 0.485 is the mean and the 0.229 is the standard deviation. In all honesty, I got those numbers from ChatGPT.

```
transform = transforms.Compose([
    transforms.Grayscale(),
    transforms.ToTensor(), # Convert image to tensor
    transforms.Normalize(0.485, 0.229) # Normalize image
])
```



Tying everything together

```
#the transformation we will apply to the images from the FER2013 dataset
transform = transforms.Compose([
    transforms.Grayscale(),
    transforms.ToTensor(), # Convert image to tensor
    transforms.Normalize(0.485, 0.229) # Normalize image
])


# loading the data from the directory I have stored the downloaded FER2013 dataset
train_data = torchvision.datasets.FER2013(root='/content', split = 'train', transform=transform)
```

Apply this series of
transformations to each
sample


Create a new Dataset
with the FER2013
blueprint

Find the dataset
at this location in
the file system

Only retrieve the
training data



```
def train_pl():  
    #the transformation we will apply to the images from the FER2013 dataset  
    transform = transforms.Compose([  
        transforms.Grayscale(),  
        transforms.ToTensor(), # Convert image to tensor  
        transforms.Normalize(0.485, 0.229) # Normalize image  
    ])  
  
    # loading the data from the directory I have stored the downloaded FER2013 dataset  
    train_data = torchvision.datasets.FER2013(root='/content', split = 'train', transform=transform)  
  
    # create dataloaders so that the FER2013 data can be loaded into the model we will implement  
    train_loader = torch.utils.data.DataLoader(train_data, batch_size=19, shuffle=True, num_workers=2)  
  
    return train_loader
```



PyTorch DataLoader

```
# create dataloaders so that the FER2013 data can be loaded into the model we will implement  
train_loader = torch.utils.data.DataLoader(train_data, batch_size=19, shuffle=True, num_workers=2)
```

There's multiple parts to this statement:

- torch.utils.data.DataLoader
- 1st parameter: train_data
- 2nd parameter: batch_size = 19
- 3rd parameter: shuffle = True
- 4th parameter: num_workers = 2

Let's look at each part




PyTorch DataLoader

DataLoader: a class in PyTorch that sets up your dataset in a way where it could be fed into your model.


It allows you to:

- Cut your data into batches
- Iterate through your data/batches
- Shuffle your data to prevent overfitting
- Create multiple pipelines
- etc.
- **More on these later...**



torch.utils.data.DataLoader

- “torch.utils.data.DataLoader” is how you access the code for DataLoader (for beginners, think of how you use System.out.print() in Java).
- “torch” is the PyTorch library
- “utils” is a specific module in the library
- “data” is another module within “utils”
- “DataLoader” is the code for the class within “data”
- Info on Python Modules can be found [here](#)



1st parameter: train_data

- The DataLoader class needs information from you first to create a new DataLoader object
 - Kind of like when you need ingredients in order to follow a cooking recipe
- The first one being the actual data.
- It needs to know the dataset that you want to train your model on.
- Only accepts PyTorch datasets.
- This is why we converted the dataset we have into a PyTorch dataset and stored it in a variable called "train_data"



2nd parameter: batch_size = 19 (**BATCH TRAINING**)

Batch Training: a machine learning technique that is commonly used for training deep learning models.

- We cut down the dataset into multiple batches
 - In this case, multiple batches with 19 data points.
- We feed the deep learning those smaller batches instead of the whole dataset at once
- More efficient because we aren't overloading the model
- More accurate because we're having the model train more.

What's the catch?

- It requires more memory on your computer, but since we are using Google Colab, we don't need to worry about that (as much)

See proceeding slides for example and illustration



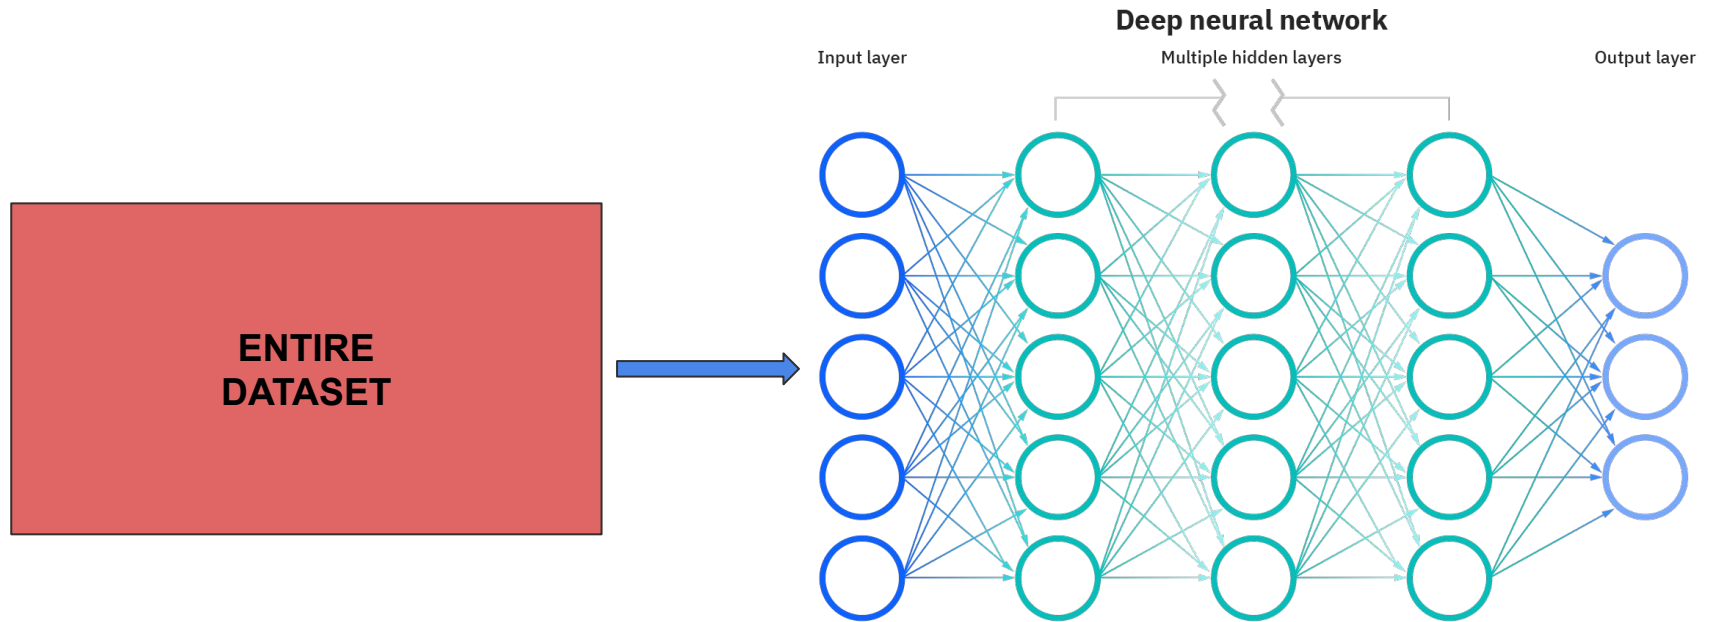
Motivating Example for Batch Training



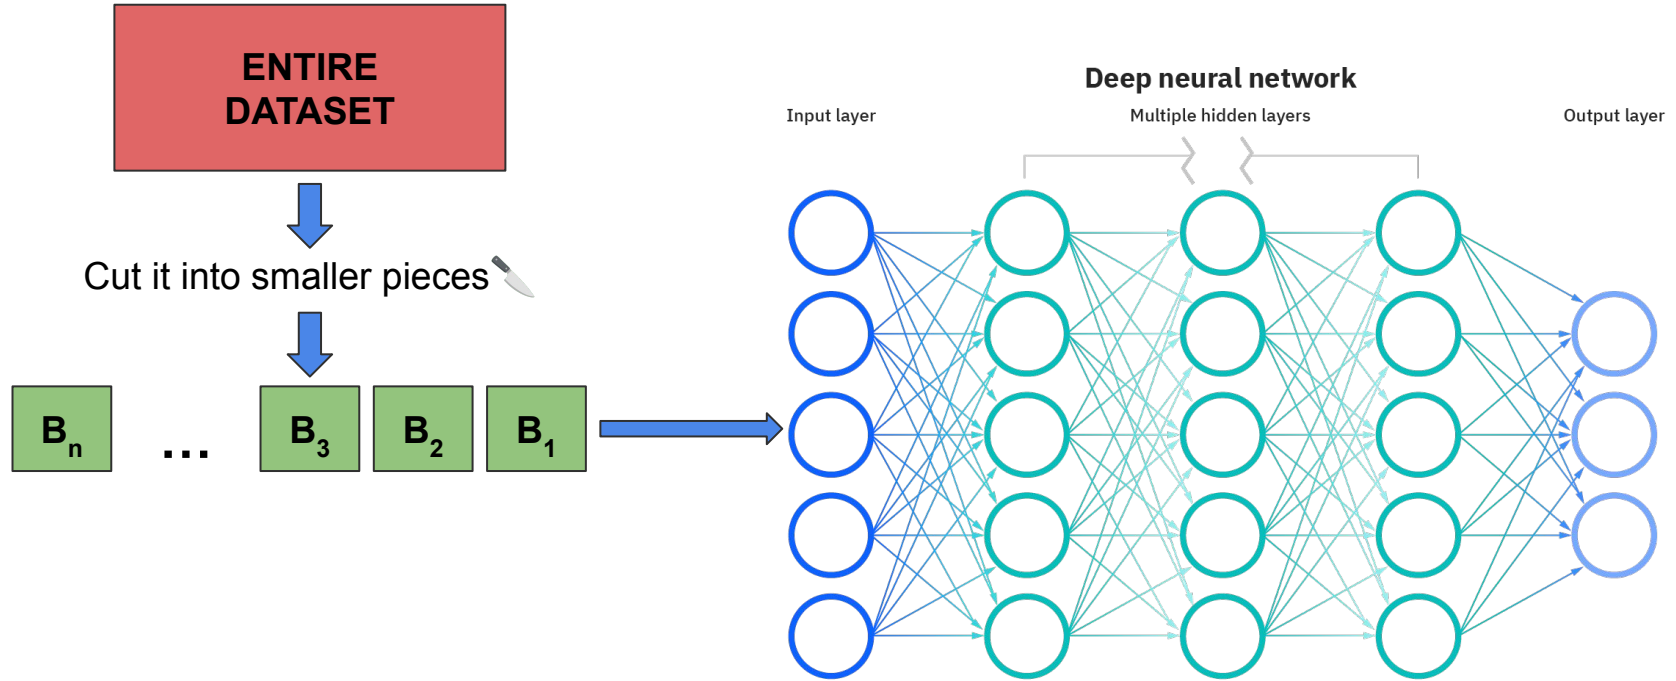
Consider eating a large piece of steak:

- Would you eat the steak whole?
- Or cut it into smaller pieces?
- Eating the steak whole would take a lot of time, effort, strain, and you don't get to enjoy it as much.
- Cutting it into pieces is easier, more efficient, and you get to enjoy your food.

No Batch Training



vs. Batch Training



3rd parameter: shuffle = True

- We're setting the shuffle parameter to true.
- We're taking the batches we made and mixing them up.
- We want to do this to prevent something called **overfitting**
- **Overfitting**: when you deep learning models memorizes the data it trained on
 - It'll do really well everytime you train, but it'll do really bad when you actually test it.
 - Imagine if you were trying to teach a kid how to read. Would you have them memorize a paragraph, or try to understand the words in the paragraph?



4th parameter: num_workers = 2

- “num_workers” is a parameter that allows you to create worker processes.
- These worker processes branch off from your main process.
- Your main process is how the data is being loaded in by default.
- Allows for parallel data loading.
 - Higher speeds when training your model
 - Requires more memory on your computer
- If you only use the main process, all of your data will be fed to the model in a single feed.
- If you use the main process + worker processes, your data will be split up fed to the model in multiple feeds.
- See next slide for example



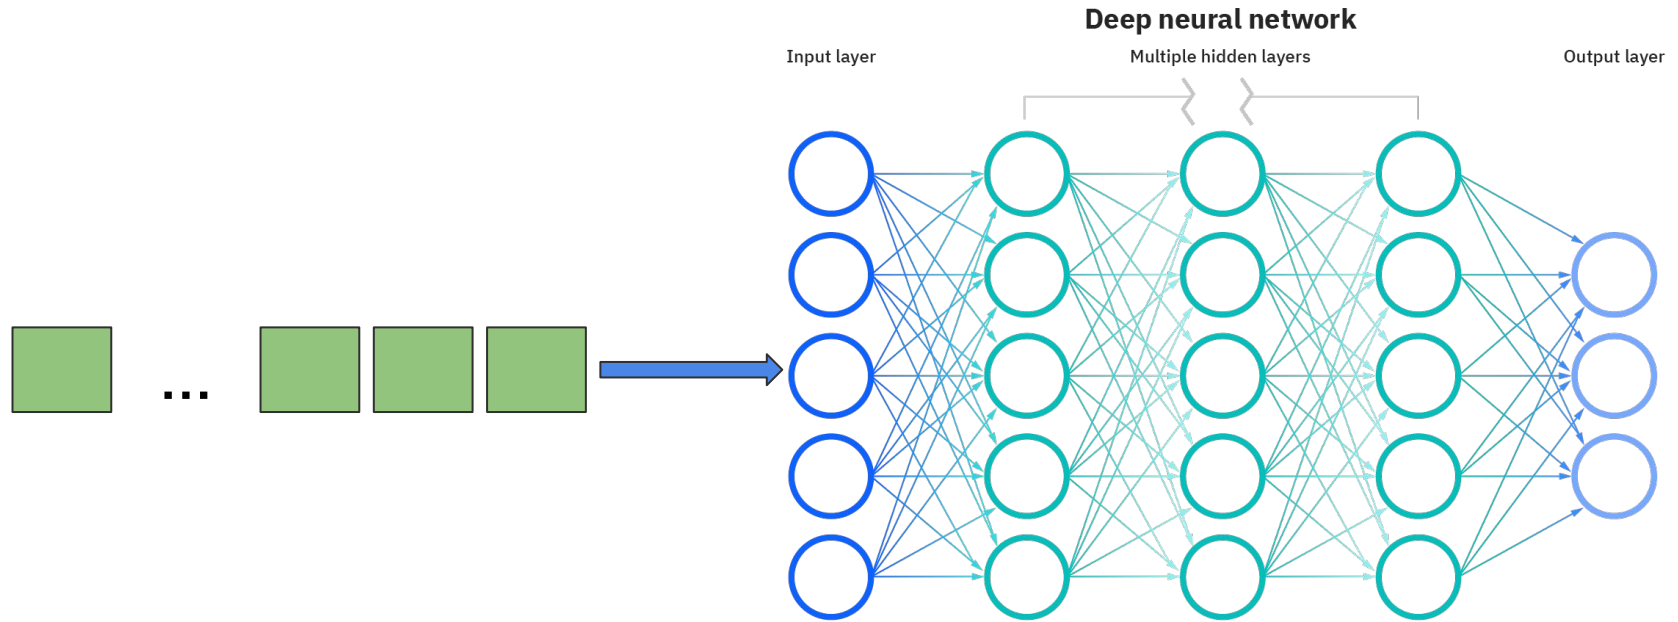
Motivating Example for num_workers



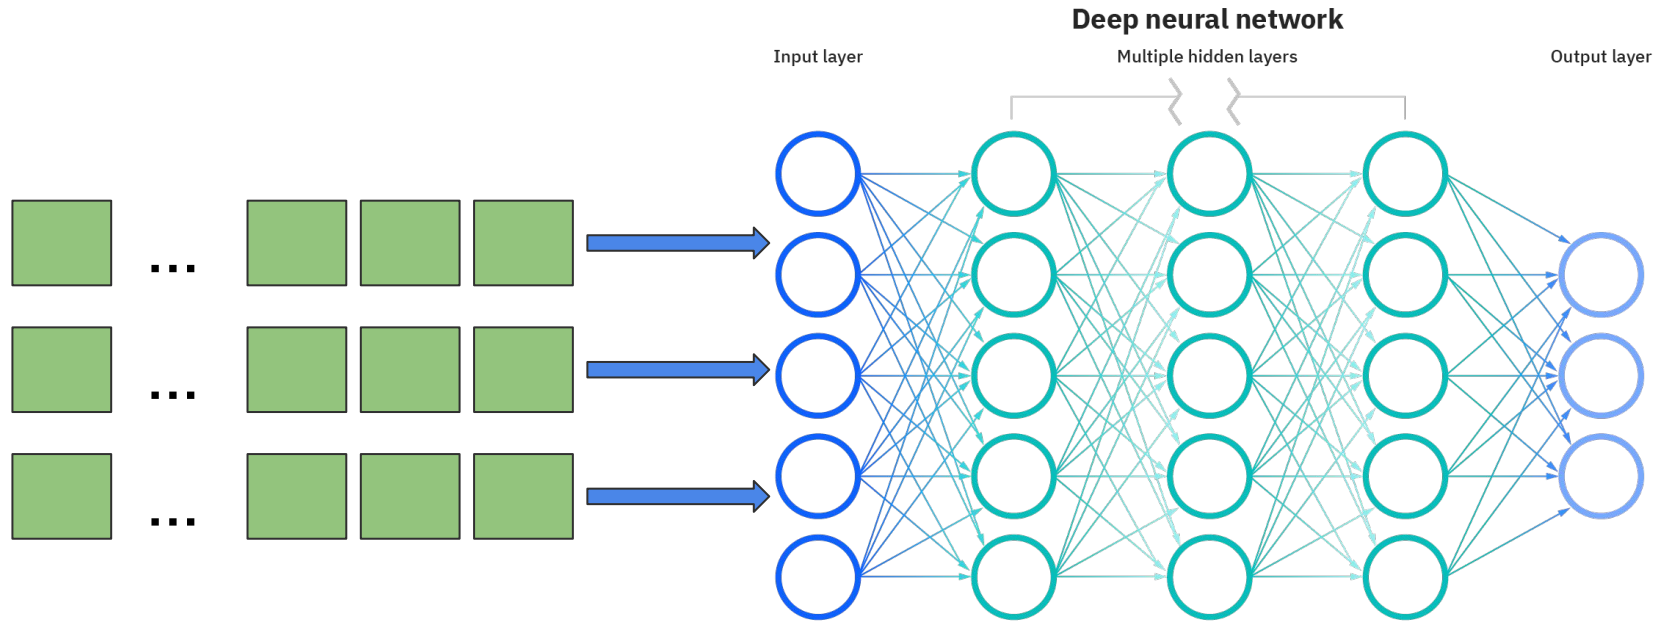
Consider TSA at airports

- A lot of times there's way too many people and security only has 1-3 checkpoints open for people to pass through
- Inefficient, inaccurate, and a real pain in the ass
- Wouldn't it be better to have more checkpoints open?

Main process only



vs. main process + worker processes



Tying everything together

```
# create dataloaders so that the FER2013 data can be loaded into the model we will implement  
train_loader = torch.utils.data.DataLoader(train_data, batch_size=19, shuffle=True, num_workers=2)
```

Create a new
DataLoader
object

Using the
FER2013 train
dataset

Split the train
data up into
batches of size
19

Shuffle the
batches

Create 2 extra
worker
processes (main
+ worker +
worker = 3
streams)

Congrats!

You now have your dataset fully set up and ready to go

- With this, you will be able to feed your data into the model you will make.

```
def train_pl():  
    #the transformation we will apply to the images from the FER2013 dataset  
    transform = transforms.Compose([  
        transforms.Grayscale(),  
        transforms.ToTensor(), # Convert image to tensor  
        transforms.Normalize(0.485, 0.229) # Normalize image  
    ])  
  
    # loading the data from the directory I have stored the downloaded FER2013 dataset  
    train_data = torchvision.datasets.FER2013(root='/content', split = 'train', transform=transform)  
  
    # create dataloaders so that the FER2013 data can be loaded into the model we will implement  
    train_loader = torch.utils.data.DataLoader(train_data, batch_size=19, shuffle=True, num_workers=2)  
  
    return train_loader
```

Next week: creating a model



What “making a model”
sounds like



What it's actually like
(with PyTorch)

From now on meetings are in Hitchcock 031!!

Will be posted in Announcements, Discord, and changed on Website





Enjoy your week!

go.osu.edu/aiclubsignup

