



Autumn 2023 Week 5

go.osu.edu/aiclubsignup



IBM Z Career Connection The Ohio State University

**Learn about IBM Z: the extremely powerful computers
making the financial world go round with a focus on
security, AI, and open source tech!**

Pomerene Hall Rm 160

October 16

6:00-8:00pm ET

ibm.biz/OSU-ZCC

Sponsored by



Sponsored by



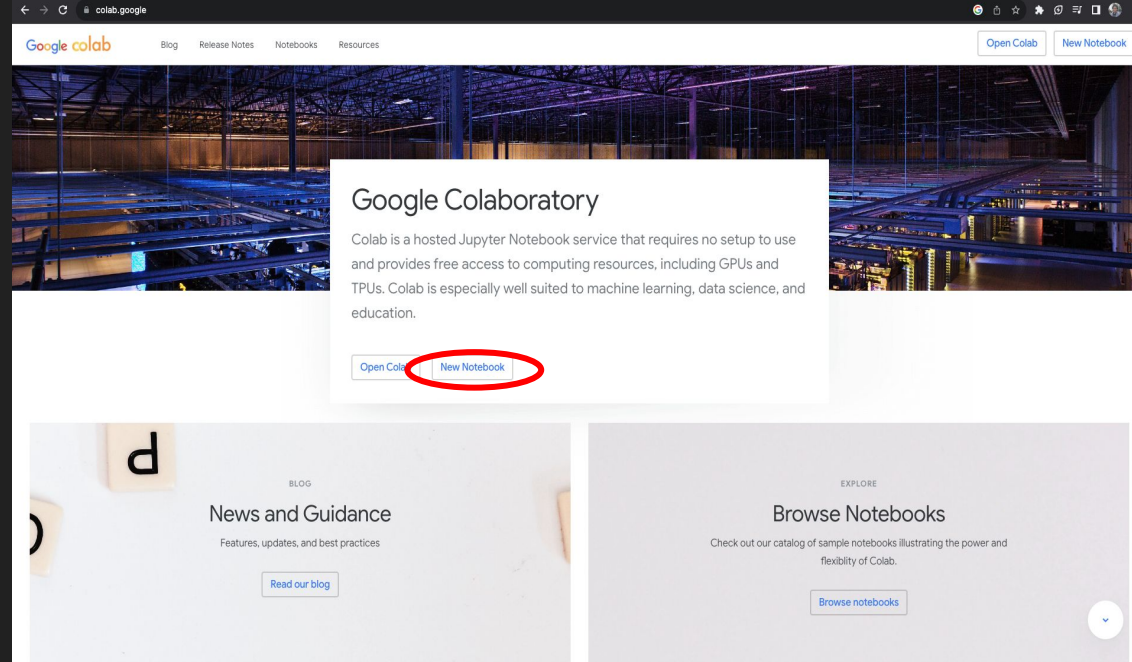
Plan for today

- What goes behind training a model
- How to train your model



Before we start...

- Open up Google Colab
- Open your previous notebook
- <https://colab.research.google.com/>



Enter this into your import code block

```
from torch import nn, save, load
from torch.optim import Adam
```



ProjectSeriesLiveLessonipynb ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

Getting the dataset into your google colab

```
[ ] !unzip /content/challenges-in-representation-learning-facial-expression-recognition-challenge.zip
```

Importing necessary libraries

```
[ ] import torch
import torchvision
from torchvision import transforms
```

Data Preprocessing & Loading


```
#this is for data preprocessing and loading with train data
def train_pl():
    #the transformation we will apply to the images from the FER2013 dataset
    transform = transforms.Compose([
        transforms.Grayscale(),
        transforms.ToTensor(), # Convert image to tensor
        transforms.Normalize(0.485, 0.229) # Normalize image
    ])

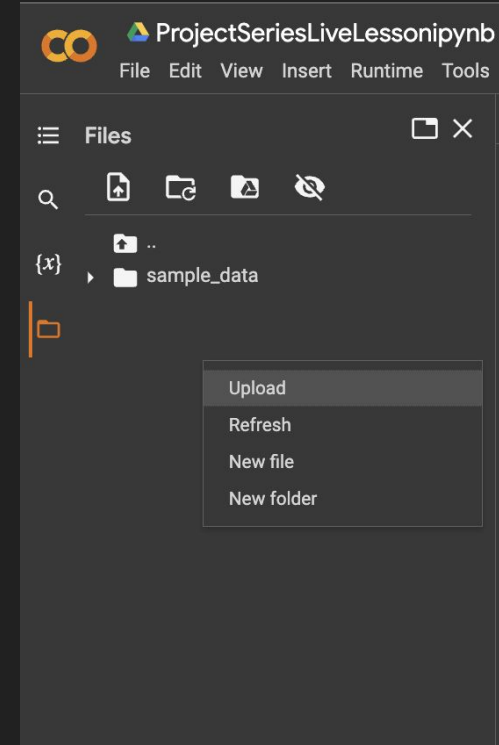
    # loading the data from the directory I have stored the downloaded FER2013 dataset
    train_data = torchvision.datasets.FER2013(root='/content', split = 'train', transform=transform)

    # create dataloaders so that the FER2013 data can be loaded into the model we will implement
    train_loader = torch.utils.data.DataLoader(train_data, batch_size=19, shuffle=True, num_workers=2)

    return train_loader
```

Downloading Dataset into Google Colab (renewed)

- Click on the  icon on the left side bar.
- Then right click in the file area and click “Upload” (as shown in the picture to the right)
- Then upload the zip file that we downloaded in the previous slide



Downloading Dataset into Google Colab (renewed)


- Create a new code block with the button that says “ + Code ” at the top.
- Type in this line into that code block and click run (if your zip file is called something else, replace the “challenges-in-representation-learning-facial-expression-recognition-challenge” part with the name of your zip file
- Then wait a minute or two for the files to show up on your colab files.
- DELETE this code block from your notebook once you’ve done this!!!



```
!unzip /content/challenges-in-representation-learning-facial-expression-recognition-challenge.zip
```


Downloading Dataset into Google Colab (renewed)

- Once all the files unzipped, right click in the file area and click “New folder” (as shown in the picture to the right)
- Title this folder “fer2013” (MAKE SURE TO COPY PASTE THIS EXACT)



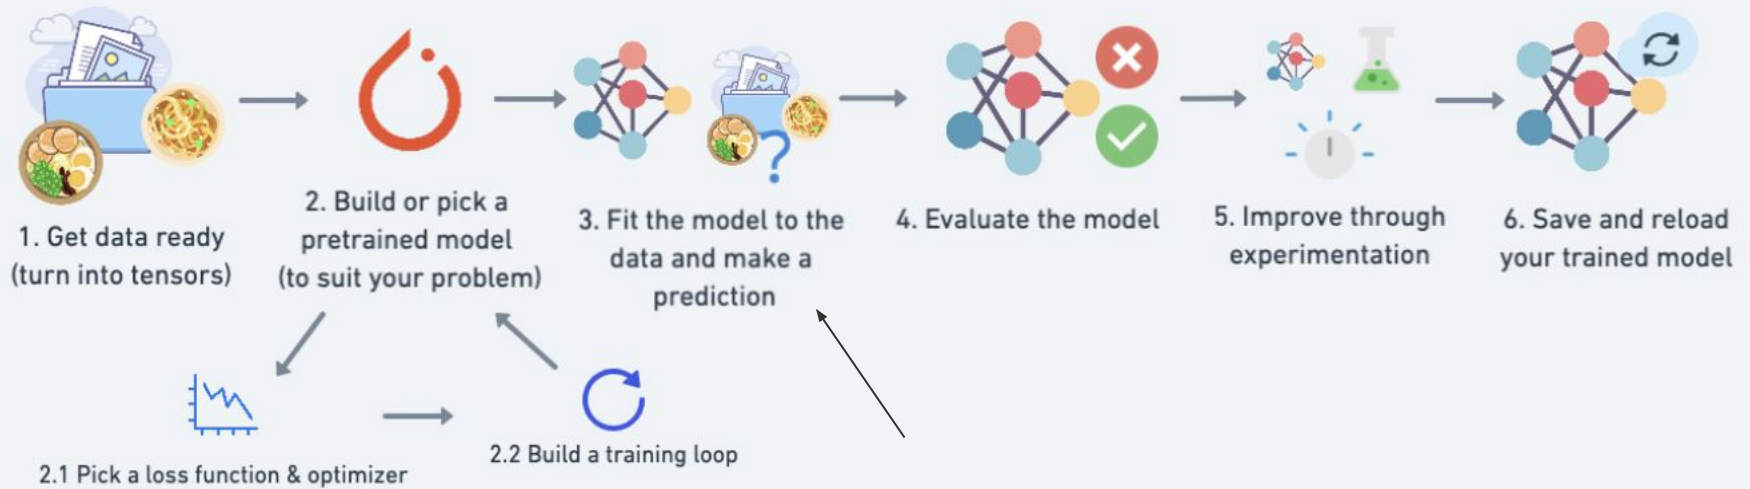
Downloading Dataset into Google Colab (renewed)

- Move the “train.csv” and “test.csv” files into the “fer2013” folder you just made.
- You should have something that looks like the picture to the right.



Our Goal:

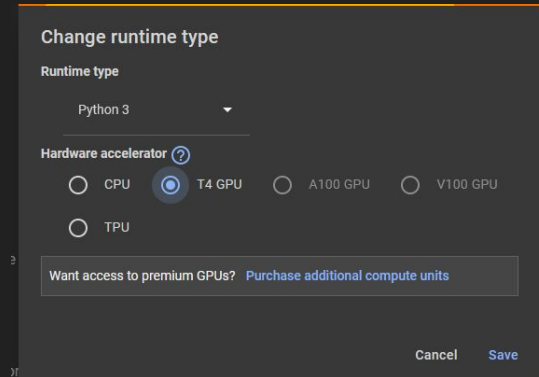
A PyTorch Workflow



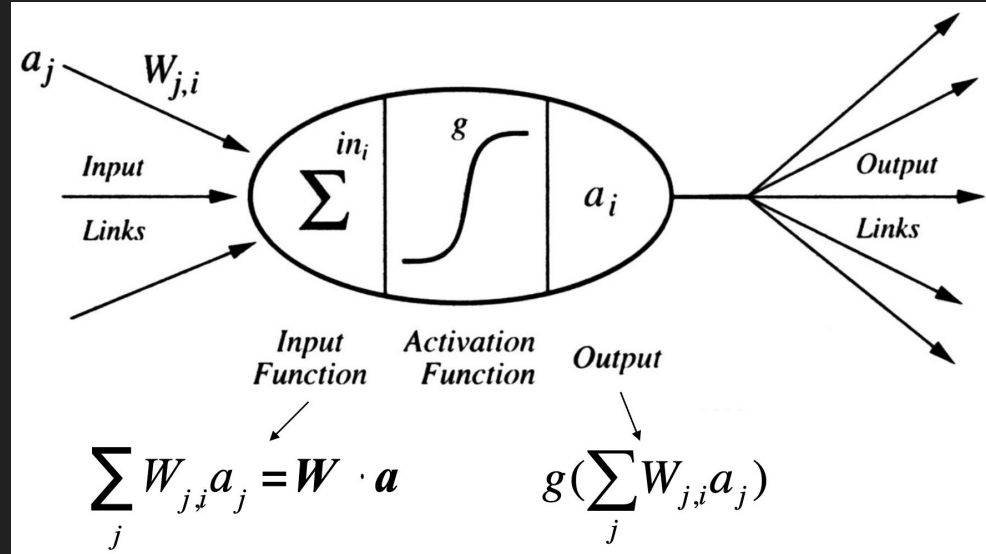
Setting Colab up with a GPU:

WITH the runtime prepared (all data set up):

1. Click the “Runtime” button under the notebook title
2. Click “Change runtime type” in the dropdown
3. Select T4 GPU and click “Save”



Context: The Anatomy of a Neuron



Some Definitions:

- **Loss:** A numeric value representing how bad your model is.
- **Loss Function:** The equation used to calculate the loss.
- **Optimizer:** The algorithm used to adjust the model and decrease the *Loss* after each batch of data.
- **Epoch:** One run of the training loop, AKA one training session using all of the training data.



Loss and the Loss Function

Cool depictions of Loss Functions: <https://losslandscape.com/gallery/>

In our case (image classification), the ideal loss function is “Cross Entropy Loss”.

The theory behind cross entropy loss is worthy of its own lecture...

[v7Labs Overview of Cross Entropy](#)

To calculate model loss, the outputs are compared to the expected outputs.

The closer the loss is to 0, the better the model is.



Gradient Descent

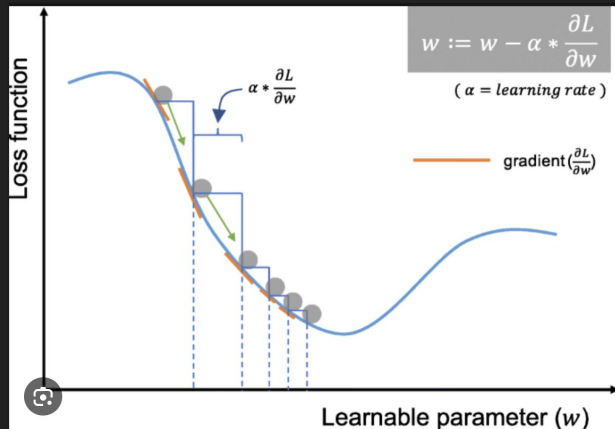
Definition: The process of gradually minimizing the loss value

Think about it like descending down a hill.

Calculus connection: Finding the local minima of a 2D function.

In our model, this process will be done with the Adam optimizer, an improvement on stochastic gradient descent. Here's an overview:

Stochastic Gradient Descent



What does training actually do?

During one epoch, the model will...

1. Do a “forward pass”, of each batch of training data
 - a. Loss is calculated after each batch
2. Calculate the average loss of the model
3. Update the parameters of each neuron
 - a. Parameters = weights - think back to the anatomy of a neuron
 - b. This process is called backpropagation (again, theory worthy of a whole lecture)

Our code for training the model

```
[ ] def train():  
    model = EmotionModel().to('cuda')  
    optimizer = Adam(model.parameters(), lr = 1e-3)  
    loss_fn = nn.CrossEntropyLoss()  
  
    train_set = train_pl()  
  
    for epoch in range(50): #train for 50 epochs  
        for batch in train_set:  
            X, y = batch  
            X, y = X.to('cuda'), y.to('cuda')  
            prediction = model(X)  
            loss = loss_fn(prediction, y)  
  
            optimizer.zero_grad()  
            loss.backward()  
            optimizer.step()  
  
            print(f"Epoch {epoch+1}\n-----")  
            print(f"\tloss:{loss}")  
            print("-----")  
  
    # saving our model to our environment  
  
    return model
```

1. Setting up the needed components and variables: the model, an optimizer, a loss function, and the training data set (more on optimizer and loss fn later)
2. Training the model over and over again

Component and Variable Initializations

```
model = EmotionModel().to('cuda')
optimizer = Adam(model.parameters(), lr = 1e-3)
loss_fn = nn.CrossEntropyLoss()


train_set = train_pl()
```

- First line creates a new instance of our model and sets it up on a GPU (faster to train on GPU than CPU).
- Sets up an optimizer called Adam.
- Sets up a cross entropy loss function.
- Sets up the training dataset.
- More on the optimizer and loss function later.

First Time Sign up



AI Club Website





Enjoy your week!

go.osu.edu/aiclubsignup

